

## PDF examples: Composite

```
CompositeThrustFunctor::CompositeThrustFunctor
(std::string n, FunctorBase* core, FunctorBase* shell)
: ThrustPdfFunctor(0, n)
{
    std::vector<unsigned int> pindices;
    pindices.push_back(core->getFunctionIndex());
    pindices.push_back(core->getParameterIndex());
    pindices.push_back(shell->getFunctionIndex());
    pindices.push_back(shell->getParameterIndex());

    // Add as components so that observables and
    // parameters will be registered.
    components.push_back(core);
    components.push_back(shell);

    cudaMemcpyFromSymbol((void**) &host_fcn_ptr,
                        ptr_to_Composite, sizeof(void*));
    initialise(pindices);
}
```

```
__device__ fptype device_Composite (fptype* evt,
                                   fptype* p,
                                   unsigned int* indices) {
    unsigned int coreFcnIndex  = indices[1];
    unsigned int coreParIndex  = indices[2];
    unsigned int shellFcnIndex = indices[3];
    unsigned int shellParIndex = indices[4];

    // NB, not normalising core function, it is not being used as a PDF.
    fptype coreValue = callFunction(evt, coreFcnIndex, coreParIndex);

    unsigned int* shellParams = paramIndices + shellParIndex;
    unsigned int numShellPars = shellParams[0];
    unsigned int shellObsIndex = shellParams[2 + numShellPars];

    fptype fakeEvt[10]; // Allow plenty of space in case events are large.
    fakeEvt[shellObsIndex] = coreValue;

    // Don't normalise shell either, since we don't know what composite
    // function is being used for. It may not be a PDF. Normalising at
    // this stage would be presumptuous.
    fptype ret = callFunction(fakeEvt, shellFcnIndex, shellParIndex);
    return ret;
}
```

## GooFit global variables

- `__constant__ fptype cudaArray[maxParams]`: Holds fit parameters - in effect Variables.
- `__constant__ unsigned int paramIndices[maxParams]`: Holds the parameter indices.
- `__constant__ fptype functorConstants[maxParams]`: Non-integer constants. Note: The zeroth entry is reserved for the number of events.
- `__constant__ fptype normalisationFactors[maxParams]`: Holds the *inverse* normalisation integrals of the PDFs.
- `__device__ void* device_function_table[200]`: Function lookup table.
- `fptype* cudaDataArray`: Stores events.
- `int host_callnumber = 0`: MINUIT iteration - host side.
- `int totalParams = 0`: Registered parameters.
- `int totalConstants = 1`: Registered constants.
- `cpuDebug, gpuDebug, debugParamIndex` - explained below.

## PDF example: Products

```
ProdThrustFunctor::ProdThrustFunctor
(std::string n, std::vector<FunctorBase*> comps)
: ThrustPdfFunctor(0, n)
, varOverlaps(false)
{
    std::vector<unsigned int> pindices;

    for (std::vector<FunctorBase*>::iterator p = comps.begin();
        p != comps.end(); ++p) {
        assert(*p);
        components.push_back(*p);
    }

    getObservables(observables); // Gathers from components

    FunctorBase::obsCont observableCheck;
    // Use to check for overlap in observables

    // Indices stores
    // (function index)(function parameter index)(variable index)
    // for each component.
    for (std::vector<FunctorBase*>::iterator p = comps.begin();
        p != comps.end(); ++p) {
        pindices.push_back((*p)->getFunctionIndex());
        pindices.push_back((*p)->getParameterIndex());
    }
}
```

```

    if (varOverlaps) continue; // Only need to establish this once.
    FunctorBase::obsCont currObses;
    (*p)->getObservables(currObses);
    for (FunctorBase::obsIter o = currObses.begin(); o != currObses.end(); ++o) {
        if (find(observableCheck.begin(), observableCheck.end(), (*o)) ==
            observableCheck.end()) continue;
        varOverlaps = true;
        break;
    }
    (*p)->getObservables(observableCheck);
}

if (varOverlaps) { // Check for components forcing separate normalisation
    for (std::vector<FunctorBase*>::iterator p = comps.begin();
        p != comps.end(); ++p) {
        if ((*p)->getSpecialMask() & FunctorBase::ForceSeparateNorm)
            varOverlaps = false;
    }
}

cudaMemcpyFromSymbol((void**) &host_fcn_ptr, ptr_to_ProdPdfs, sizeof(void*));
initialise(pindices);
}

```

```

__host__ fptype ProdThrustFunctor::normalise () const {
    if (varOverlaps) {
        // Two or more components share an observable
        // and cannot be separately normalised, since
        // \int A*B dx does not equal \int A dx * \int B dx.
        recursiveSetNormalisation(fptype(1.0));
        cudaMemcpyToSymbol(normalisationFactors,
                           host_normalisation,
                           totalParams*sizeof(fptype),
                           0, cudaMemcpyHostToDevice);
        // Normalise numerically.
        fptype ret = ThrustPdfFunctor::normalise();
        return ret;
    }

    // Normalise components individually
    for (std::vector<FunctorBase*>::const_iterator c = components.begin();
         c != components.end(); ++c) {
        (*c)->normalise();
    }
    host_normalisation[parameters] = 1;
    cudaMemcpyToSymbol(normalisationFactors,
                       host_normalisation,
                       totalParams*sizeof(fptype),
                       0, cudaMemcpyHostToDevice);

    return 1.0;
}

```

```
__device__ fptype device_ProdPdfs (fptype* evt,
                                   fptype* p,
                                   unsigned int* indices) {
    // Index structure is nP | F1 P1 | F2 P2 | ...
    // where nP is number of parameters, Fs are
    // function indices, and Ps are parameter indices

    int numParams = indices[0];
    fptype ret = 1;
    for (int i = 1; i < numParams; i += 2) {
        int fcnIdx = indices[i + 0];
        int parIdx = indices[i + 1];

        fptype curr = callFunction(evt, fcnIdx, parIdx);
        curr *= normalisationFactors[parIdx];
        ret *= curr;
    }

    return ret;
}
```

---

## Member variables of FunctorBase

- **parameters:** Host-side parameter index for each PDF.
- **cIndex:** Host-side constants index - points into functorConstants.
- **normRanges:** Array storing lower limit, upper limit and step size of normalisation integral.
- **observables:** List of independent variables this PDF depends on. Superset of all component observables.
- **parameterList:** List of parameters of this PDF.
- **components:** List of ‘nested’ PDFs, eg the factors of a product.
- **cachedParams:** Parameter values from previous MINUIT iteration - used to check if, for example, normalisation integrals need to be recalculated.



## Tricks for debugging

- If you look at the original (unsanitised) versions of the GooFit PDFs, they are full of commented-out statements like this:

```
/*  
    if (0 == blockIdx.x + threadIdx.x)  
        printf("Conv debug: %i %i %i %i %i %i\n",  
            numOthers,  
            indices[8],  
            numToLoad,  
            workspaceIndex,  
            offsetInBins,  
            numbins);  
*/
```

- Clearly, their development is not always entirely smooth...
- This is very nearly my only trick!
- ThrustPdfFunctor has a `setDebugMask(int, bool)` method, which I occasionally use to switch particular debug statements on and off, eg:

```
if (gpuDebug & 1) printf("Whatever\n");
```

The method sets global variables `cpuDebug` and `gpuDebug`. Additionally, if its boolean parameter is true, the variable `debugParamIndex` is set to equal the parameter index of the PDF object on which the method was called. I use

---

this to print debug statements only from a specific PDF in cases when I have several of the same type in my fit; for example:

```
if ((gpuDebug & 1) && (paramIndices + debugParamIndex == indices))  
    printf("Whatever\n");
```

- Sometimes it is helpful to look at error codes; raw CUDA calls generally return an error, eg:

```
cudaError_t err = cudaMemcpy(target, source, size, cudaMemcpyHostToDevice);  
if (cudaSuccess != err) {  
    std::cout << "Error near line " << __LINE__  
              << " : " << cudaGetErrorString(err)  
              << std::endl;  
    abort(0);  
}
```

- It can in principle also be helpful to use `cuda-gdb`, which works much like ordinary `gdb`. Compile with `-G` instead of (or in addition to) `-g`.

## PDF examples: Time-dependent mixing


- This is a very specialised PDF.
- Start with merely describing the quasi-two-body resonances that make up the total decay amplitude:

```
struct ResonanceInfo {
    Variable* amp_real;
    Variable* amp_imag;
    Variable* mass;
    Variable* width;
    unsigned int spin;
    unsigned int cyclic_index;
    unsigned int eval_type;
};
struct DecayInfo {
    fptype motherMass;
    fptype daug1Mass;
    fptype daug2Mass;
    fptype daug3Mass;
    fptype meson_radius;
    Variable* _tau;
    Variable* _xmixing;
    Variable* _ymixing;
    std::vector<ResonanceInfo*> resonances;
};
```

- The core of the function is a sum over Breit-Wigner resonances:

$$\mathcal{A}(m_1, m_2) = \sum_{j=0}^N a_j e^{i\theta_j} B_j(m_1, m_2; m_j, \sigma_j)$$

- Amplitudes  $a_j e^{i\theta_j}$  vary quickly; resonance masses and widths ( $m_j, \sigma_j$ ) are kept fixed. Breit-Wigners are relatively costly to evaluate.
- Obviously: Cache the Breit-Wigner values at every point in the Dalitz plot!



## Create the cache

- Now we need to know the data size so the special *wave calculator* knows how far to jump, given that it's on event N.

```
__host__ void TddpThrustFunctor::setDataSize (unsigned int dataSize,
                                              unsigned int evtSize) {
    // Default 5 is m12, m13, time, sigma_t, evtNum
    totalEventSize = evtSize;
    assert(totalEventSize >= 5);

    if (cachedWaves) {
        delete cachedWaves;
    }

    numEntries = dataSize;
    cachedWaves = new device_vector<WaveHolder>
        (dataSize*decayInfo->resonances.size());
    void* dummy = thrust::raw_pointer_cast(cachedWaves->data());
    cudaMemcpyToSymbol(cWaves, &dummy, sizeof(WaveHolder*),
        cacheToUse*sizeof(WaveHolder*));
    setForceIntegrals();
}
```

## Special normalisation

```
for (int i = 0; i < decayInfo->resonances.size(); ++i) {
    if (!redoIntegral[i]) continue;

    strided_range<device_vector<WaveHolder>::iterator> aragorn
        (cachedWaves->begin() + i,
         cachedWaves->end(),
         decayInfo->resonances.size());

    transform(make_zip_iterator(make_tuple(eventIndex, dataArray, eventSize)),
              make_zip_iterator(make_tuple(eventIndex + numEntries,
                                             arrayAddress, eventSize)),
              aragorn.begin(),
              *(calculators[i]));
}
// End of time-consuming integrals.
```

- Input iterator is just looping over indices.
- Output iterator does a *strided* loop over `cachedWaves`, jumping by the number of resonances every time input iterator jumps by one.
- Also cache integrals of each pair of resonances.

---

## Summary

- We can write *very complex* PDFs and still get good speedups.
- Effort to complexity ratio is similar to RooFit's.
- Speed is way better!
- Practice makes perfect.