
Plan for the day

- Engine internals
- Timing tools
- Hardware
- TODOs
- Discussion, wrapup

Engine internals

- Three levels of GooFit code:
 - User level: Create `FooThrustFunctor` and `DataSet` objects and use in a fit.
 - Development: Write new PDF classes.
 - Framework: The core engine that ties together the classes.
- Engine code is mainly concentrated in `ThrustPdfFunctor.cu`, with some outliers in `FunctorBase.cu` and `PdfBuilder.cc`.
- Following the program flow, let's begin with `PdfBuilder.cc`, containing the class `PdfFunctor`.

```
set<Variable*> vars;  
FunctorBase* pdfPointer = 0;  
  
void PdfFunctor::fit () {  
    host_callnumber = 0;  
    vars.clear();  
    pdfPointer->getParameters(vars);  
  
    numPars = vars.size();  
    if (minuit) delete minuit;  
    minuit = new TMinuit(numPars);  
    int maxIndex = 0;  
    int counter = 0;  
    for (set<Variable*>::iterator i = vars.begin(); i != vars.end(); ++i) {
```

```

minuit->DefineParameter(counter, (*i)->name.c_str(),
                        (*i)->value, (*i)->error,
                        (*i)->lowerlimit, (*i)->upperlimit);
if ((*i)->fixed) minuit->FixParameter(counter);
counter++;
if (maxIndex < (*i)->getIndex()) maxIndex = (*i)->getIndex();
}

numPars = maxIndex+1;
pdfPointer->copyParams();

minuit->SetFCN(FitFun);
if (0 < overrideCallLimit) {
    std::cout << "Calling MIGRAD with call limit "
                << overrideCallLimit << std::endl;
    double plist[1];
    plist[0] = overrideCallLimit;
    int err = 0;
    minuit->mnexcm("MIGRAD", plist, 1, err);
}
else minuit->Migrad();
}

void FitFun(int &npar, double *gin, double &fun, double *fp, int iflag) {
    vector<double> pars;
    // Notice that npar is number of variable parameters, not total.
    pars.resize(numPars);

```

```
int counter = 0;

for (set<Variable*>::iterator i = vars.begin(); i != vars.end(); ++i) {
    pars[(*i)->getIndex()] = fp[counter++];
}

pdfPointer->copyParams(pars);
fun = pdfPointer->calculateNLL();
host_callnumber++;
}
```

calculateNLL function

```
__host__ double ThrustPdfFunctor::calculateNLL () const {
    normalise();
    cudaMemcpyToSymbol(normalisationFactors,
                        host_normalisation,
                        totalParams*sizeof(fptype),
                        0, cudaMemcpyHostToDevice);
    cudaDeviceSynchronize(); // Ensure normalisation integrals are finished

    int numVars = observables.size();
    if (fitControl->binnedFit()) {
        numVars += 2;
        numVars *= -1;
    }

    fptype ret = sumOfNll(numVars);
    if (0 == ret)
        abortWithCudaPrintFlush(__FILE__, __LINE__, getName() + " zero NLL", this);
    return 2*ret;
}

__host__ double ThrustPdfFunctor::sumOfNll (int numVars) const {
    static plus<double> cudaPlus;
    constant_iterator<int> eventSize(numVars);
    constant_iterator<fptype*> arrayAddress(cudaDataArray);
    double dummy = 0;
```

```
counting_iterator<int> eventIndex(0);  
return transform_reduce(make_zip_iterator(  
    make_tuple(eventIndex,  
                arrayAddress, eventSize)),  
    make_zip_iterator(  
        make_tuple(eventIndex + numEntries,  
                    arrayAddress, eventSize)),  
    *logger, dummy, cudaPlus);  
}
```

MetricTaker functor

- Created by the initialise method:

```
__host__ void ThrustPdfFunctor::initialise
(std::vector<unsigned int> pindices, void* dev_functionPtr) {
    if (!fitControl) setFitControl(new UnbinnedNllFit());

    // MetricTaker must be created after FunctorBase
    // initialisation is done.
    FunctorBase::initialiseIndices(pindices);

    functionIdx = findFunctionIdx(dev_functionPtr);
    // Actually a setMetric call
    logger = new MetricTaker(this, getMetricPointer(fitControl->getMetric()));
}
```

- Constructor takes a ThrustPdfFunctor (to extract the function index) and a goodness-of-fit pointer:

```
MetricTaker::MetricTaker (FunctorBase* dat, void* dev_functionPtr)
: metricIndex(0)
, functionIdx(dat->getFunctionIndex())
, parameters(dat->getParameterIndex())
{

    std::map<void*, int>::iterator localPos =
        functionAddressToDeviceIndexMap.find(dev_functionPtr);
```

```

    if (localPos != functionAddressToDeviceIndexMap.end()) {
        metricIndex = (*localPos).second;
    }
    else {
        metricIndex = num_device_functions;
        host_function_table[num_device_functions] = dev_functionPtr;
        functionAddressToDeviceIndexMap[dev_functionPtr] = num_device_functions;
        num_device_functions++;
        cutilSafeCall(cudaMemcpyToSymbol(device_function_table,
                                          host_function_table,
                                          num_device_functions*sizeof(void*)));
    }
}

```

- Operator method calls the device-side function:

```

// Notice that operators are distinguished by the order of the operands,
// and not otherwise! It's up to the user to make his tuples correctly.
// Main operator: Calls the PDF to get a predicted value, then the metric
// to get the goodness-of-prediction number which is returned to MINUIT.
__device__ fptype MetricTaker::operator () (tuple<int, fptype*, int> t) const
    int eventIndex = thrust::get<0>(t);
    int eventSize  = thrust::get<2>(t);
    fptype* eventAddress = thrust::get<1>(t) + (eventIndex * abs(eventSize));

// Causes stack size to be statically undeterminable.
fptype ret = callFunction(eventAddress, functionIdx, parameters);

```



```

// Notice assumption here! For unbinned fits the 'eventAddress' pointer
// won't be used in the metric, so it doesn't matter what it is. For
// binned fits it is assumed that the structure of the event is
// (obs1 obs2... binentry binvolume),
// so that the array passed to the metric consists of (binentry binvolume);
// unless the data has user-provided errors, in which case binvolume is
// replaced by binError.

device_metric_ptr metricFcn =
    reinterpret_cast<device_metric_ptr>(device_function_table[metricIndex]);
ret = (*metricFcn)(ret, eventAddress + (abs(eventSize)-2), parameters);
return ret;
}

// Operator for binned evaluation, no metric.
// Used in normalisation.
__device__ fptype MetricTaker::operator () (tuple<int, int, fptype*> t) const
// Bin index, event size, base address [lower, upper, numbins]

int evtSize = thrust::get<1>(t);
assert(evtSize <= MAX_NUM_OBSERVABLES);
int binNumber = thrust::get<0>(t);
__shared__ fptype binCenters[1024*MAX_NUM_OBSERVABLES];

```

```

// To convert global bin number to (x,y,z...) coordinates: For each
// dimension, take the mod with the number of bins in that dimension.
// Then divide by the number of bins, in effect collapsing so the
// grid has one fewer dimension. Rinse and repeat.

unsigned int* indices = paramIndices + parameters;
for (int i = 0; i < evtSize; ++i) {
    fptype lowerBound = thrust::get<2>(t)[3*i+0];
    fptype upperBound = thrust::get<2>(t)[3*i+1];
    int numBins      = (int) FLOOR(thrust::get<2>(t)[3*i+2] + 0.5);
    int localBin = binNumber % numBins;

    fptype x = upperBound - lowerBound;
    x /= numBins;
    x *= (localBin + 0.5);
    x += lowerBound;
    binCenters[indices[indices[0]+2+i]+threadIdx.x*MAX_NUM_OBSERVABLES] = x;
    binNumber /= numBins;
}

// Causes stack size to be statically undeterminable.
fptype ret = callFunction(binCenters+threadIdx.x*MAX_NUM_OBSERVABLES,
                          cudaArray, parameters);
return ret;
}

```

Default numerical normalisation

```
__host__ fptype ThrustPdfFunctor::normalise () const {
    if (!fitControl->metricIsPdf()) {
        host_normalisation[parameters] = 1.0;
        return 1.0;
    }

    fptype ret = 1;
    if (hasAnalyticIntegral()) {
        for (obsConstIter v = obsCBegin(); v != obsCEnd(); ++v) {
            // Loop goes only over observables of this PDF.
            ret *= integrate((*v)->lowerlimit, (*v)->upperlimit);
        }
        host_normalisation[parameters] = 1.0/ret;
        return ret;
    }

    int totalBins = 1;
    for (obsConstIter v = obsCBegin(); v != obsCEnd(); ++v) {
        ret *= ((*v)->upperlimit - (*v)->lowerlimit);
        totalBins *= (integrationBins > 0 ? integrationBins : (*v)->numbins);
    }
    ret /= totalBins;

    fptype dummy = 0;
    static plus<fptype> cudaPlus;
    constant_iterator<fptype*> arrayAddress(normRanges);
```

```
constant_iterator<int> eventSize(observables.size());
counting_iterator<int> binIndex(0);

fptype sum = transform_reduce(make_zip_iterator(
    make_tuple(binIndex,
                eventSize, arrayAddress)),
    make_zip_iterator(
    make_tuple(binIndex + totalBins,
                eventSize, arrayAddress)),
    *logger, dummy, cudaPlus);

if (isnan(sum)) {
    abortWithCudaPrintFlush(__FILE__, __LINE__,
        getName() + " NaN in normalisation", this);
}
ret *= sum;

if (0 == ret) abortWithCudaPrintFlush(__FILE__, __LINE__, "Zero integral");
host_normalisation[parameters] = 1.0/ret;
return (fptype) ret;
}
```

Timing tools

- Old standby: Look at the clock!

```
gettimeofday(&startTime, NULL);  
fitter.fit();  
gettimeofday(&stopTime, NULL);
```

- For measuring device time more strictly, we have the cudaEvent class:

```
cudaEvent_t start;  
cudaEvent_t stops;  
cudaEventCreate(&start, 0); // Int is stream index  
cudaEventCreate(&stops, 0);  
  
cudaEventRecord(start, 0); // Place start event in execution queue  
some_kernel<<<blocks, threads>>>(args);  
cudaEventRecord(stops, 0);  
cudaEventSynchronize(stop); // Now safe to read results.  
  
float timeUsed = 0;  
cudaEventElapsedTime(&timeUsed, start, stop);
```

- nVidia offers some profiling tools.
- We can also compile with profiling info; first set some environment variables:

```
export COMPUTE_PROFILE=1  
export COMPUTE_PROFILE_LOG=cuda_profile_%d.csv
```

```
export COMPUTE_PROFILE_CSV=1
export COMPUTE_PROFILE_CONFIG=./prof_config
```

and create a profiler configuration:

```
gridsize
threadblocksize
memtransfersize
memtransferdir
regperthread
dynsmemperblock
stasmemperblock
divergent_branch
local_load
local_store
```

which tells the profiler which information to store about the kernel. Compile with debug flags:

```
nvcc -G -g -o myBinary myCode.cu
```

- Running the executable will now result in a csv file that you can load into Excel or otherwise mess with:

kernel	cpu time
_Z6kernelPh,	911852.312, 911894.000,1000,1000,
	gpu time grid size

```
registers
13,          0.167,
            occupancy
```

- Can also run in visual mode. Login with -X option:

```
ssh -X username@oakley.osc.edu
qsub -X -I -l nodes=1:ppn=6:gpus=1 -l walltime=01:00:00
module load cuda
nvvp&
```

- Produces visual timeline and some helpful hints.
- Finally, latest versions of GooFit have inbuilt profiling option:

```
#ifdef PROFILING
__constant__ fptype conversion = (1.0 / CLOCKS_PER_SEC);
__device__ fptype callFunction (fptype* eventAddress, unsigned int functionIdx,
    clock_t start = clock();
    fptype ret = (*(reinterpret_cast<device_function_ptr>(device_function_table[
es + paramIdx));
    clock_t stop = clock();
    if ((0 == threadIdx.x + blockIdx.x) && (stop > start)) {
        // Avoid issue when stop overflows and start doesn't.
        timeHistogram[functionIdx*100 + paramIdx] += ((stop - start) * conversion)
    }
    return ret;
}
```

```
#else  
// (Regular callFunction method)  
#endif
```